

Lecture 4: Time Series Examples

Francis J. DiTraglia

March 19, 2014

1 Some Time Series Examples

So far we've looked at some completely generic examples (AIC, TIC, cross-validation) and some regression examples (Mallows, AIC corrected). The generic examples can immediately be applied to an ML problem, including time series setting: if you have an ML routine, you already get everything you need for these criteria as a side effect. Could be Kalman filter or conditional likelihood.

We derived the other examples (Mallows, AIC corrected) for a regression problem, but it's easy to adapt these to AR and VAR models. As long as we're willing to use conditional ML (drop some observations) these already *are* regression problems.

We'll take a look at AR and VAR models using conditional likelihood so we can write out explicit formulas for the criteria. We won't do TIC since we can't really unpack the penalty term. We'll treat cross-validation in its own section.

The consistent criteria are ... and the efficient criteria are ...

We won't go through all of the specifics of the derivations for mallows and AICc since they're almost identical to the regression derivation. Some more details can be found in McQuarrie and Tsai (1998).

1.1 Autoregressive Models

For simplicity assume there is no constant term. Then the $\text{AR}(p)$ model is

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \epsilon_t$$

where $\epsilon_t \sim \text{iid } N(0, \sigma^2)$ and we observe a sample y_1, \dots, y_N . We'll use conditional maximum likelihood, so we lose the first p observations. Thus the *effective sample size* is $T = N - p$. The conditional ML estimator of $\boldsymbol{\phi} = (\phi_1, \dots, \phi_p)'$ is simply the least-squares estimator

$$\hat{\boldsymbol{\phi}} = (X'X)^{-1}X'\mathbf{y}$$

where $\mathbf{y} = (y_{p+1}, y_{p+2}, \dots, y_N)'$ and the design matrix is

$$X = \begin{bmatrix} y_p & y_{p-1} & \cdots & y_1 \\ y_{p+1} & y_p & \cdots & y_2 \\ \vdots & \vdots & & \vdots \\ y_{N-1} & y_{N-2} & \cdots & y_{N-p-1} \end{bmatrix}$$

The maximum likelihood estimator of σ^2 is

$$\hat{\sigma}_p^2 = \frac{\text{RSS}_p}{T}$$

where RSS denotes the residual sum of squares, namely $\|\mathbf{y} - X\hat{\boldsymbol{\phi}}\|^2$. Since this is a regression model, it's trivial to adapt both Mallows's C_p and the AIC_C to this case.¹ For Mallows's C_p we have

$$C_p = \frac{\text{RSS}_p}{\hat{\sigma}_{\text{wide}}^2} - T + 2p$$

¹If you'd like to see all of the details written out, consult McQuarrie & Tsai (1998), Chapter 3.

where $\hat{\sigma}_{wide}^2$ is the estimate of σ^2 from the model with *maximum order* among those under consideration. For AIC_c we have

$$AIC_c = \log(\hat{\sigma}_p^2) + \frac{T + p}{T - p - 2}$$

For both C_p and AIC_c we choose the lag length that *minimizes* the criterion.

Using an argument essentially identical to the one presented in the notes for Lecture 2, the maximized log-likelihood for the $AR(p)$ model is

$$-\frac{T}{2} [\log(2\pi) + \log(\hat{\sigma}_p^2) + 1]$$

To construct the AIC and BIC, we multiply this quantity by 2 and subtract the appropriate penalty term, ignoring terms that are constant across models. The number of parameters for an $AR(p)$ model is $p + 1$, since we estimate σ^2 in addition to the p autoregressive parameters. We'll rescale both AIC and BIC and flip their signs to make them comparable to the C_p and AIC_c expressions from above. Putting everything together for the sake of comparison, we have

$$\begin{aligned} AIC &= \log(\hat{\sigma}_p^2) + \frac{2(p + 1)}{T} \\ AIC_c &= \log(\hat{\sigma}_p^2) + \frac{T + p}{T - p - 2} \\ C_p &= \frac{RSS_p}{\hat{\sigma}_{wide}^2} + 2p - T \\ BIC &= \log(\hat{\sigma}_p^2) + \frac{\log(T)(p + 1)}{T} \end{aligned}$$

In each case, we choose the model that *minimizes* the criterion. Of these four criteria, only BIC is consistent. The other three criteria, however, are efficient under one-step-ahead squared prediction error loss in an environment in which the true DGP is an infinite-order autoregression. The BIC does not have this property.

Ng & Perron (2005) There are some subtle but important points that we glossed over in the preceding discussion and that are, indeed, rarely mentioned in textbooks or articles on model selection. First there is the question of whether we should use the maximum likelihood estimator $\hat{\sigma}^2$ or the unbiased estimator that divides by $T - p$ rather than T . In time series applications T may be small enough that it makes a difference. More troubling, however, is the problem of deciding what should count as the sample size, since different lag lengths use a different number of observations in the conditional maximum likelihood setting. Indeed, as they are usually written, expressions for AIC and BIC drop terms that are constant across models in *cross-section regression*, where changing the number of regressors doesn't affect sample size. The situation is of course entirely different for AR models but practitioners *still use the same formulas* in this case. There are numerous different ways to handle these complications. Ng & Perron (2005) review the possibilities and illustrate how each performs in a number of simulation studies.

1.2 Vector Autoregression Models

Again, assume the intercept is zero. Then the VAR(p) model is given by

$$\begin{aligned} \mathbf{y}_t &= \Phi_1 \mathbf{y}_{t-1} + \dots + \Phi_p \mathbf{y}_{t-p} + \boldsymbol{\epsilon}_t \\ \mathbf{y}_t &_{(q \times 1)} \quad \quad \quad \Phi_i &_{(q \times q)} \\ \boldsymbol{\epsilon}_t &\stackrel{iid}{\sim} N_q(\mathbf{0}, \Sigma) \end{aligned}$$

where we observe $\mathbf{y}_1, \dots, \mathbf{y}_N$. Again, if we're content to use conditional maximum likelihood, dropping the first p observations to estimate a VAR(p) model, this is simply a multivariate regression problem and we have an *effective sample size* of $T = N - p$. Written as a multivariate regression model, we have

$$\mathbf{Y} = \mathbf{X} \boldsymbol{\Phi} + \mathbf{U}$$

$(T \times q) \quad (T \times pq)(pq \times q) \quad (T \times q)$

where

$$Y_{(T \times q)} = \begin{bmatrix} \mathbf{y}'_{p+1} \\ \mathbf{y}'_{p+2} \\ \vdots \\ \mathbf{y}'_N \end{bmatrix}, \quad \Phi_{(pq \times q)} = \begin{bmatrix} \Phi'_1 \\ \Phi'_2 \\ \vdots \\ \Phi'_p \end{bmatrix}, \quad U_{(T \times q)} = \begin{bmatrix} \epsilon'_{p+1} \\ \epsilon'_{p+2} \\ \vdots \\ \epsilon'_N \end{bmatrix}$$

and the design matrix is

$$X_{(T \times pq)} = \begin{bmatrix} \mathbf{y}'_p & \mathbf{y}'_{p-1} & \cdots & \mathbf{y}'_1 \\ \mathbf{y}'_{p+1} & \mathbf{y}'_p & \cdots & \mathbf{y}'_2 \\ \vdots & \vdots & & \vdots \\ \mathbf{y}'_{N-1} & \mathbf{y}'_{N-2} & \cdots & \mathbf{y}'_{N-p-1} \end{bmatrix}$$

Thus, the conditional maximum likelihood estimator for Φ is

$$\hat{\Phi} = (X'X)^{-1}X'Y$$

and the maximum likelihood estimator for Σ is

$$\hat{\Sigma}_p = \frac{(Y - X\hat{\Phi})'(Y - X\hat{\Phi})}{T}$$

The VAR(p) model has a very large number of parameters. First, we have the coefficients of Φ_1, \dots, Φ_p . Each of these is an unrestricted $q \times q$ matrix so Φ contains a total of pq^2 parameters. We also need to estimate the variance matrix Σ of the errors ϵ . Although Σ contains q^2 elements, it is a symmetric matrix so there are only $q(q+1)/2$ free parameters. Thus, a VAR(p) model requires us to estimate a total of $pq^2 + (q+1)q/2$ parameters. To calculate the AIC and BIC we also need the maximized log-likelihood, which is given by

$$-\frac{T}{2} \left[q \log(2\pi) + \log |\hat{\Sigma}_p| + q \right]$$

Re-scaling as we did for the AR model, we have

$$\text{AIC} = \log \left| \widehat{\Sigma}_p \right| + \frac{2pq^2 + q(q+1)}{T}$$

$$\text{BIC} = \log \left| \widehat{\Sigma}_p \right| + \frac{\log(T)pq^2}{T}$$

The multivariate generalization of AIC_c is

$$\text{AIC}_c = \log \left| \widehat{\Sigma}_p \right| + \frac{(T + qp)q}{T - qp - q - 1}$$

as explained in Chapter 5 of McQuarrie and Tsai (1998). For each of the preceding three expressions, we choose the model that *minimizes* the given criterion. Of these criteria, both AIC and its corrected version are efficient while BIC is consistent.

1.3 Corrected AIC for State Space Models

As the lag length p grows, the number of parameters in a $\text{VAR}(p)$ model explodes, and can easily come close to the effective sample size. In situations like this, AIC is known to perform poorly. The bias correction $2 \times \text{length}(\theta)$ is based on a large-sample argument and fails to provide a good approximation when the number of parameters is too close to the sample size, leading the AIC to choose models that are in general “too large” to achieve our target of minimizing the KL divergence.² The idea behind the AIC_c of Hurvich and Tsai (1989) was to provide a better approximation to the AIC bias correction for AR models under a certain set of assumptions. In a similar vein, Cavanaugh & Shumway (1997) propose a refined AIC, the AIC_b , for general state space models. Rather than deriving an analytical correction term, they suggest using the bootstrap to approximate the bias of the maximized log-likelihood

²Cavanaugh & Shumway (1997) suggest $\text{length}(\theta) \approx T/2$ as a rough approximation of what counts as “too many parameters relative to sample size” for the AIC to work well.

as an estimator of the expected log likelihood, using the state-space bootstrap procedure proposed by Stoffer and Wall (1991).

2 More on Cross-Validation

How to extend it to time series. Varieties other than leave-one-out. Efficiency versus consistency. Racine (2000) and Burman, Chow & Nolan (1994).

2.1 K-Fold Cross-validation

In Lecture 2 we discussed leave-one-out (LOO) cross-validation (CV) and proved that it is asymptotically equivalent to TIC. In fact LOO-CV can be viewed as a special case of a more general procedure: **K-fold cross-validation**. The idea behind cross-validation is to create a “stand-in” for the – typically infeasible – ideal of fitting your model on a “training” dataset and assessing its predictive performance out-of-sample on an independent “validation” dataset. In the simplest and most common case we suppose that we have observed an iid sample and K-fold CV proceeds as follows

1. Randomly partition the dataset into K parts of roughly equal size. These are the “folds.”
2. For each $k = 1, \dots, K$ do the following:
 - (a) Estimate your model using all observations *except* those contained in the k th fold.
 - (b) Using the results, construct $\left\{ \hat{y}_t^{(-k)} \right\}_{t=1}^T$ denote the predictions for y_t based on estimation *without* the observations from the k th fold.
- 3.
4. Repeat this procedure for each model under consideration, and select the one that yields the smallest value of $CV(K)$.

How to choose K ? Bias-Variance tradeoff. Correlation between estimators. Computational considerations: K -fold is easier than LOO. Often use $K = 5$ or 10 . For linear smoothers and quadratic prediction loss, don't need to actually do the LOO computations. You proved a special case of this on the most recent problem set.

2.2 How to handle dependent observations

AR example.

Cross-Validation for AR The way we described it above, CV depended in independence. How can we adapt it for AR models? Roughly speaking, the idea is to use the fact that dependence dies out over time and treat observations that are “far enough apart” as *approximately* independent. Specifically, we choose an integer value h and assume that y_t and y_s can be treated as independent as long as $|s - t| > h$. This idea is called “ h -block cross-validation” and was introduced by Burman, Chow & Nolan (1994). As in the iid version of leave-one-out cross-validation, we still evaluate a loss function by predicting *one* withheld observation at a time using a model estimated without it. The difference is that we also omit the h neighboring observations *on each side* when fitting the model. For example, if we choose to evaluate squared-error loss, the criterion is

$$CV_h(1) = \frac{1}{T - p} \sum_{t=p+1}^T (y_t - \hat{y}_{(t)}^h)^2$$

where

$$\hat{y}_{(t)}^h = \hat{\phi}_{1(t)}^h y_{t-1} + \dots + \hat{\phi}_{1(t)}^h y_{t-p}$$

and $\hat{\phi}_{j(t)}^h$ denotes the j th parameter estimate from the conditional least-squares estimator with observations y_{t-h}, \dots, y_{t+h} removed. We still have the question of what h to choose. Here there is a trade-off between making the assumption

of independence more plausible and leaving enough observations to get precise model estimates. Intriguingly, the simulation evidence presented in McQuarrie and Tsai (1998) suggests that setting $h = 0$, which yields plain-vanilla leave-one-out CV, works well even in settings with dependence.

The idea of h -block cross-validation can also be adapted to versions of cross-validation other than leave-one-out. For details, see Racine (1997, 2000).

Cross-Validation for VARs In principle we could use the same h -block idea here as we did for the AR example above. However, given the large number of parameters we need to estimate, the sample sizes withholding $2h + 1$ observations at a time may be too small for this to work well.

2.3 Consistent or Efficient?

Shao (1993) and Racine (1997, 2000)

3 Two Additional Criteria

We've already covered the most commonly used model selection criteria, but there are two others that come up from time to time: Akaike's Final Prediction Error (FPE), and the Hannan-Quinn Information Criterion (HQ). Roughly speaking FPE behaves like AIC while HQ behaves like BIC: while FPE is efficient, HQ is consistent.

3.1 Final Prediction Error

3.2 Hannan-Quinn

In our last lecture we examined a consistency result based on the central limit theorem. It is also possible to construct a consistency proof by appealing to the law of the iterated logarithm. This is how HQ is derived.

For VAR models

$$FPE = \left| \widehat{\Sigma}_p \right| \left(\frac{T + qp}{T - qp} \right)^q$$

$$HQ = \log \left| \widehat{\Sigma}_p \right| + \frac{c \log \log(T) pq^2}{T}$$

where $c > 2$. A commonly-used value is 2.01.